

# Assignment\_1

Yuying Cheng

```
```{r setup}
#| message: false
#| warning: false

library(ggplot2)
library(knitr)
library(kableExtra)
library(dplyr)
library(tidyverse)
library(stringr)
library(RColorBrewer)

fixed_palette <- "Dark2"
set_style <- function(p){
  return(p +
    theme_classic() +
    theme(legend.position = "top") +
    scale_color_brewer(palette = fixed_palette)+
    scale_fill_brewer(palette = fixed_palette))
}
```
```

## Task 4 - R basic operations

```
sqrt(10)
```

```
[1] 3.162278
```

```
log2(32)
```

```
[1] 5
```

```
sum <- 0
for (i in seq(1, 1000)){
  sum <- sum + i
}
sum
```

```
[1] 500500
```

```
sum <- 0
for (i in seq(2, 1000, 2)){
  sum <- sum + i
}
sum
```

```
[1] 250500
```

```
choose(100, 2)
```

```
[1] 4950
```

```
choose(100, 3)
```

```
[1] 161700
```

## Task 5 - Using R example datasets

**Describe briefly the content of the CO2 dataset using the help function.**

CO2 is data frame contains data from an experiment on cold tolerance of the grass species *Echinochloa crus-galli*. The experiment subjects are 6 plants originated either from Quebec or Mississippi. Their CO<sub>2</sub> uptake rate was measured at several levels of ambient CO<sub>2</sub>, with 2 treatment conditions - chilled and not chilled before the measurement.

**What is the average and median CO2 uptake of the plants from Quebec and Mississippi?**

```
data(CO2)
head(CO2) %>%
  kbl() %>%
  kable_styling()
CO2 %>%
  group_by(Type) %>%
  summarise(median = median(uptake), average = mean(uptake)) %>%
  kbl() %>%
  kable_styling()
```

| Plant | Type   | Treatment  | conc | uptake |
|-------|--------|------------|------|--------|
| Qn1   | Quebec | nonchilled | 95   | 16.0   |
| Qn1   | Quebec | nonchilled | 175  | 30.4   |
| Qn1   | Quebec | nonchilled | 250  | 34.8   |
| Qn1   | Quebec | nonchilled | 350  | 37.2   |
| Qn1   | Quebec | nonchilled | 500  | 35.3   |
| Qn1   | Quebec | nonchilled | 675  | 39.2   |


| Type        | median | average  |
|-------------|--------|----------|
| Quebec      | 37.15  | 33.54286 |
| Mississippi | 19.30  | 20.88333 |

**In the “airway” example data from Bioconductor, how many genes are expressed in each sample? How many genes are not expressed in any sample?**

```
library(airway)
data(airway)
expressed_no <- list(
  expressed = ~sum(.x > 0),
  not_expressed = ~sum(.x == 0)
)
df <- as.data.frame(assay(airway)) %>%
  summarise_all(.fun = expressed_no,
    .names = "{.col}.{.fn}") %>%
  gather(sample, count) %>%
  mutate(category = str_extract(sample, "expressed|not_expressed"),
    sample = str_remove(sample, ".expressed|.not_expressed"))
df %>%
  kbl() %>%
  kable_styling()
```

```
p <- ggplot(df, aes(x = sample, y = count, fill = category)) +
  geom_bar(stat = "identity")
set_style(p) +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1))
```

| sample     | count | category      |
|------------|-------|---------------|
| SRR1039508 | 24633 | expressed     |
| SRR1039509 | 24527 | expressed     |
| SRR1039512 | 25699 | expressed     |
| SRR1039513 | 23124 | expressed     |
| SRR1039516 | 25508 | expressed     |
| SRR1039517 | 25998 | expressed     |
| SRR1039520 | 24662 | expressed     |
| SRR1039521 | 23991 | expressed     |
| SRR1039508 | 39469 | not_expressed |
| SRR1039509 | 39575 | not_expressed |
| SRR1039512 | 38403 | not_expressed |
| SRR1039513 | 40978 | not_expressed |
| SRR1039516 | 38594 | not_expressed |
| SRR1039517 | 38104 | not_expressed |
| SRR1039520 | 39440 | not_expressed |
| SRR1039521 | 40111 | not_expressed |




## Task 6 - R Functions

Write a function that calculates the ratio of the mean and the median of a given vector.

```
#Calculates the ratio of mean and median of a given vector
#input: a numeric vector
#output: the ratio of mean and median of the vector
r_mean_median <- function(v){
  return(mean(v)/median(v))
}

#test
v <- list(normal = rnorm(10000, mean = 1),
          exponential = rexp(10000, rate = 1))
hist(v$normal, breaks = 100)
```

```
hist(v$exponential, breaks = 100)
lapply(v, r_mean_median)
```



```
$normal
[1] 1.011004
```

```
$exponential
[1] 1.446122
```

**Write a function that ignores the lowest and the highest value from a given vector and calculate the mean.**

```
#Calculates mean of a vectore after removing *one* maximum and *one* minimum
#input: a vector
#output: mean value after removal of one maximum and one minimum
adjusted_average <- function(v){
  x = sum(v) - min(v) - max(v)
  return(x/(length(v)-2))
}

#test
v = c(seq(1,3), rep(10,3))
v
```

```
[1] 1 2 3 10 10 10
```

```
adjusted_average(v)
```

```
[1] 6.25
```

## Pipes

Pipe is a tool predominantly for a linear sequence of operations. To use it, connection the operations with “%>%”. However because of its design, there are several situations when it is not appropriate to use pipe, including:

1. when the number of operations is too large, piping makes it hard to debug.
2. when there are multiple inputs and outputs.
3. when there's a non-linear dependency structure of the operations.

## Apply family

The apply family is designed to replace the use of loops when it fits. In my work, the apply family is handy when i have several data sets but all subjected to the same pre-processing processes. For example, when i want to get the read counts from bigwig files, instead of writing loops I can lump the sets into a list and use `lapply` so I transform each bigwig file into a read count matrix with one line of code.

## Task 7 - Basic visualization



**Compare the distributions of the body heights of the two species from the ‘magic\_guys.csv’ dataset graphically.**

```
df <- read.csv("magic_guys.csv")
unique(df$species)
```

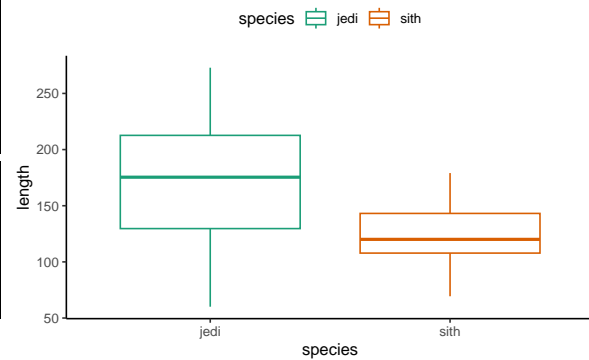

```
[1] "jedi" "sith"
```

```
hist(df[df$species == "jedi", "length"], breaks = 20, xlab = "length", main = "jedi")
hist(df[df$species == "sith", "length"], breaks = 20, xlab = "length", main = "sith")
```

```
p1 <- ggplot(data = df, aes(x = length, fill = species, color = species)) +
  geom_histogram(binwidth = 10, alpha = 0.8) +
  facet_grid(species ~ .)
p1 <- set_style(p1) + theme(legend.position = "none")
p1
p2 <- ggplot(data = df, aes(x = species, y = length, color = species)) +
  geom_boxplot()
```



```
p2 <- set_style(p2)
p2
```



```
plots <- list(p1, p2)
names(plots) <- c("gg_hist.png", "gg_box.png")
lapply(names(plots), function(x) ggsave(x, plots[[x]]))
```

Saving 5.5 x 3.5 in image  
Saving 5.5 x 3.5 in image

```
[[1]]
[1] "gg_hist.png"
```

```
[[2]]
[1] "gg_box.png"
```

Load the gene expression data matrix from the 'microarray\_data.tab' dataset provided in the shared folder, it is a big tabular separated matrix.


1. How big is the matrix in terms of rows and columns?

```
df <- read.table("microarray_data.tab", sep = "\t", header = T)
size_sum(df)
```

```
[1] "[553 x 1,000]"
```

2. Count the missing values per gene and visualize this result.

```
gene_na <- df %>%
  summarise_all(~sum(is.na(.))) %>%
  gather(key = "gene", value = "count")
p1 <- ggplot(gene_na, aes(x = gene, y = count)) +
  geom_point(alpha = 0.4) +
  ylab("NA count") +
  theme(axis.text.x = element_blank(),
        axis.ticks.x = element_blank())
p1
p2 <- ggplot(gene_na, aes(x = count)) +
  geom_density()
set_style(p2)
```



3. Find the genes for which there are more than X% (X=10%, 20%, 50%) missing values.

```
set_10 <- gene_na %>%
  filter(count > nrow(df) * 0.1) %>%
  mutate(category = ">10%")
```




```

set_20 <- set_10 %>%
  filter(count > nrow(df) * 0.2) %>%
  mutate(category = ">20%")
set_50 <- set_20 %>%
  filter(count > nrow(df) * 0.5) %>%
  mutate(category = ">50%")
set_na <- gene_na %>%
  filter(count == nrow(df))

p1 +
  geom_point(data = set_10, color = brewer.pal(4, fixed_palette)[1], alpha = 0.4) +
  geom_point(data = set_20, color = brewer.pal(4, fixed_palette)[2], alpha = 0.6) +
  geom_point(data = set_50, color = brewer.pal(4, fixed_palette)[3], alpha = 0.8) +
  geom_point(data = set_na, color = brewer.pal(4, fixed_palette)[4]) +
  geom_hline(yintercept = c(0.1, 0.2, 0.5, 1) * nrow(df), color = "grey40", linetype = "da

```



4. Replace the missing values by the average expression value for the particular gene. (Note: Imputing data has to be used with caution!)

```

replace_with_mean <- function(v){
  v[is.na(v)] <- mean(v, na.rm = T)
}

```

```

}

df_imputed <- df %>%
  mutate(across(where(is.numeric), ~replace_na(., mean(., na.rm = T))))
df_imputed[1:8, 1:6]

```

|   | g1          | g2        | g3         | g4          | g5          | g6         |
|---|-------------|-----------|------------|-------------|-------------|------------|
| 1 | 1.80200000  | 0.1656927 | -0.1820000 | 1.31200000  | 3.49700000  | 0.4390000  |
| 2 | 0.02547518  | 0.1656927 | 7.6930000  | -0.06731957 | 0.19300000  | -1.3830000 |
| 3 | 1.07900000  | 0.1656927 | 1.5560000  | 1.65200000  | -0.01812288 | 0.4600000  |
| 4 | 3.60700000  | 0.1656927 | 1.9140000  | -0.06731957 | 1.40000000  | 1.1090000  |
| 5 | -1.70000000 | 0.1656927 | 0.9430000  | -0.06731957 | -0.17000000 | -0.1571338 |
| 6 | 0.02547518  | 0.1656927 | 0.0430000  | -0.06731957 | 0.72900000  | -0.0890000 |
| 7 | 0.02547518  | 0.1656927 | -0.1230605 | -0.06731957 | -0.01812288 | -0.1571338 |
| 8 | 0.02547518  | 0.1656927 | -0.1230605 | -0.06731957 | -0.01812288 | -1.2970000 |

**Visualize the data in the CO2 dataset in a way that gives you a deeper understanding of the data. What do you see?**


```

getPallette <- colorRampPalette(brewer.pal(8, fixed_palette))
N_color <- length(unique(CO2$Plant))
p1<- ggplot(CO2, aes(x = Treatment, y = uptake, fill = Plant)) +
  geom_boxplot()
set_style(p1) + scale_fill_manual(values = getPallette(N_color))
p2 <- ggplot(CO2, aes(x = conc, y = uptake, color = Treatment)) +
  geom_point() +
  geom_smooth(method = "loess", fill = NA) +
  facet_grid(Type ~ .)
set_style(p2)
p3 <- ggplot(CO2, aes(x = conc, y = uptake, color = Type)) +
  geom_point() +
  geom_smooth(method = "loess", fill = NA) +
  facet_grid(Treatment ~ .)
set_style(p3)

```

From the above plots I noticed the following:

1. From the boxplot, plants from Quebec have a higher CO<sub>2</sub> uptake rate than Mississippi, whether they were chilled or not before the measurement.
2. Comparing the treatments, chilling decreases the CO<sub>2</sub> uptake rate in general.



- From the second and the third plot, we can learn the relationship between ambient CO<sub>2</sub> level (conc) and the CO<sub>2</sub> uptake rate of the plants. This relationship of the plants demonstrates a similar trend, i.e. CO<sub>2</sub> uptake rate increases as the level of conc increases, with a decreasing rate of change, and eventually reaches a plateau.
- Whether or not chilled before measurement affects (decreases the uptake rate) the plants from Mississippi more than Quebec.

## Task 8

Install the Tidybiology package, which includes the data ‘chromosome’ and ‘proteins’.


- Extract summary statistics (mean, median and maximum) for the following variables from the ‘chromosome’ data: variations, protein coding genes, and miRNAs. Utilize the tidyverse functions to make this as simply as possible.

```
library(tidybiology)
#colnames(chromosome)
chromosome %>% select(c("variations", "protein_codinggenes", "mi_rna")) %>%
  summarise_all(.funs = list(mean = ~round(mean(.x)), median = median, max = max),
    .names = "{.col}.{.fn}") %>%
  gather(key = "statistics") %>%
  kbl() %>%
  kable_styling()
```

- How does the chromosome size distribute? Plot a graph that helps to visualize this by using ggplot2 package functions.

```
p <- ggplot(chromosome, aes(x = id, y = length_mm)) +
  geom_bar(stat = "identity") +
  geom_text(aes(label = length_mm), vjust = 1.6, color = "white", size = 3)
set_style(p)
```


| statistics                 | value    |
|----------------------------|----------|
| variations_mean            | 6484572  |
| protein_codinggenes_mean   | 850      |
| mi_rna_mean                | 73       |
| variations_median          | 6172346  |
| protein_codinggenes_median | 836      |
| mi_rna_median              | 75       |
| variations_max             | 12945965 |
| protein_codinggenes_max    | 2058     |
| mi_rna_max                 | 134      |



3. Does the number of protein coding genes or miRNAs correlate with the length of the chromosome? Make two separate plots to visualize these relationships.

```
p1 <- ggplot(chromosome, aes(x = length_mm, y = protein_codinggenes)) +
  geom_smooth(method = "lm") +
  geom_point()
set_style(p1)
p2 <- ggplot(chromosome, aes(x = length_mm, y = mi_rna)) +
  geom_smooth(method = "lm") +
  geom_point()
```

```
set_style(p2)
```



4. Calculate the same summary statistics for the 'proteins' data variables length and mass. Create a meaningful visualization of the relationship between these two variables by utilizing the ggplot2 package functions. Play with the colors, theme- and other visualization parameters to create a plot that pleases you.

```
proteins %>% select(c("length", "mass")) %>%
  summarise_all(.funs = list(mean = mean, median = median, max = max),
                .names = "{.col}.{.fn}") %>%
  gather(key = "statistics") %>%
  kbl() %>%
  kable_styling()
p <- ggplot(proteins, aes(x = length, y = mass)) +
  geom_point()
set_style(p)
```

| statistics    | value            |
|---------------|------------------|
| length_mean   | 557.160254527655 |
| mass_mean     | 62061.3791483113 |
| length_median | 414              |
| mass_median   | 46140.5          |
| length_max    | 34350            |
| mass_max      | 3816030          |

